

IST687 - Music Classification Project

Team 2 - Sebastian Castro, John Fields, Courtney Smith, Jeremy Wallner

5/13/2019

Executive Summary

The purpose of this project is to analyze the Million Song Database to predict “Hot” artists and songs based on the attributes such as familiarity, artist location, loudness, terms used, etc. The analysis was done using R software on a 10,000 track subset of the data and our model was able to predict “Hot” songs with ~80% accuracy.

Table of Contents

Executive Summary Data Analysis Conclusion Final proofing

Introduction

Related Work

Thierry Bertin-Mahieux, Daniel P.W. Ellis, Brian Whitman, and Paul Lamere. The Million Song Dataset. In Proceedings of the 12th International Society for Music Information Retrieval Conference (ISMIR 2011), 2011.

Dataset

```
#New code from Courtney to change from 3 to 5 categories of artist hotness
music <- read.csv("/Users/johnfields/Library/Mobile Documents/com~apple~CloudDocs/Syracuse/IST687/Project1/million_song_database.csv")
str(music)
```

```
## 'data.frame': 9996 obs. of 36 variables:
## $ artist.hottness : num 0.402 0.417 0.343 0.454 0.402 ...
## $ artist.id : Factor w/ 3885 levels "AR009211187B989185",...: 1269 2353 2168 715 3606 ...
## $ artist.name : Factor w/ 4409 levels ":Blacks On :Blondes",...: 682 3796 3560 67 1569 ...
## $ artist_mbtags : Factor w/ 277 levels "", "0.333", "60s",...: 1 52 1 262 1 1 1 1 1 ...
## $ artist_mbtags_count : num 0 1 0 1 0 0 0 0 0 ...
## $ bars_confidence : num 0.643 0.007 0.98 0.017 0.175 0.121 0.709 0.142 0.806 0.047 ...
## $ bars_start : num 0.585 0.711 0.732 1.306 1.064 ...
## $ beats_confidence : num 0.834 1 0.98 0.809 0.883 0.438 0.709 0.234 0.44 1 ...
## $ beats_start : num 0.585 0.206 0.732 0.81 0.136 ...
## $ duration : num 219 148 177 233 210 ...
## $ end_of_fade_in : num 0.247 0.148 0.282 0 0.066 ...
## $ familiarity : num 0.582 0.631 0.487 0.63 0.651 ...
## $ key : num 1 6 8 0 2 5 1 4 4 7 ...
## $ key_confidence : num 0.736 0.169 0.643 0.751 0.092 0.635 0 0 0.717 0.053 ...
## $ latitude : num 37.2 35.1 37.2 37.2 37.2 ...
## $ location : Factor w/ 1046 levels " ", " NC", " UbA!", Minas Gerais",...: 157 584 705 ...
## $ longitude : num -63.9 -90 -63.9 -63.9 -63.9 ...
## $ loudness : num -11.2 -9.84 -9.69 -9.01 -4.5 ...
```

```
## $ mode : int 0 0 1 1 1 1 1 0 1 0 ...
## $ mode_confidence : num 0.636 0.43 0.565 0.749 0.371 0.557 0 0.16 0.652 0.473 ...
## $ release.id : int 300848 300822 514953 287650 611336 41838 25824 8876 358182 692313
## $ release.name : Factor w/ 7830 levels " Lazy Afternoon En Anglais",...: 2191 1746 3535 1
## $ similar : Factor w/ 2837 levels "AR00K8N11C8A41687B",...: 2408 2225 1145 304 2331
## $ song.hottnesss : num 0.602 NA NA NA 0.605 ...
## $ song.id : Factor w/ 9996 levels " Polovtsian Dances / Rimsky-Korsakov: Russian E
## $ start_of_fade_out : num 219 138 172 217 199 ...
## $ tatums_confidence : num 0.779 0.969 0.482 0.601 1 0.136 0.467 0.292 0.121 1 ...
## $ tatums_start : num 0.285 0.206 0.421 0.563 0.136 ...
## $ tempo : num 92.2 121.3 100.1 119.3 129.7 ...
## $ terms : Factor w/ 459 levels "", "8-bit", "acid jazz",...: 216 34 372 327 325 396
## $ terms_freq : num 1 1 1 0.989 0.887 ...
## $ time_signature : num 4 4 1 4 4 3 1 3 4 4 ...
## $ time_signature_confidence: num 0.778 0.384 0 0 0.562 0.454 0 0.408 0.487 0.878 ...
## $ title : Factor w/ 9705 levels "", " -start ID-",...: 3572 7526 481 7474 2531 828
## $ year : int 0 1969 0 1982 2007 0 0 0 1984 0 ...
## $ artist.hottnesss.label : Factor w/ 3 levels "Cold","Hot","Warm": 3 3 3 3 3 3 1 2 1 3 ...
```


```
##Artist Hotness Histogram
```

```
library(ggplot2)
```

```
ggplot(music, aes(x=artist.hottnesss)) + geom_histogram(color="black", fill="steelblue", alpha=0.5) +
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

Histogram: Artist Hotness



```
##Function to create descriptive statistics for artist hotness
```

```
descriptive_stats <- function(vector) { library(moments)
```

```
  result <- c(Mean=mean(vector),
```

```

        Median=median(vector),
        Min = min(vector),
        Max = max(vector),
        SD = sd(vector),
        Quantile = quantile(vector, probs = c(0.25,.50,0.75, 0.95)),
        Skewness = skewness(vector) )

    print(result)
}
descriptive_stats(music$artist.hotttnesss)

##           Mean           Median           Min           Max           SD
##    0.3857065    0.3807564    0.0000000    1.0825026    0.1434688
## Quantile.25% Quantile.50% Quantile.75% Quantile.95%    Skewness
##    0.3255062    0.3807564    0.4539300    0.6011861   -0.1483509

##Methodology for assigning artist hottness levels - uses quantiles from descriptive_statistics functi
##95% Quantile: 0.6011861 - Hot
##75% Quantile: 0.453858 - Warm
##50% Quantile: 0.3807423 - Tepid
##25% Quantile: 0.3252656 - Cool
##Code for assigning labels based on above quantiles
music$artist.hottness.label <- ifelse(music$artist.hotttnesss >=0.6011861, "Hot",
                                     ifelse(music$artist.hotttnesss >=0.453858 & music$artist.hotttnesss < 0.6011861, "Warm",
                                             ifelse(music$artist.hotttnesss >=0.3807423 & music$artist.hotttnesss < 0.453858, "Tepid",
                                                    ifelse(music$artist.hotttnesss >=0.3252656 & music$artist.hotttnesss < 0.3807423, "Cool",
                                                           ifelse(music$artist.hotttnesss < 0.3252656, "Frigid", "Hot"))))
unique(music$artist.hottness.label)

## [1] "Tepid" "Cool" "Warm" "Frigid" "Hot"

#End of new code from Courtney
#Prior to importing, a new column artist.hotttnesss.label was adding with
#Hot(>.4590), Warm(<.4590 and >.3357), Cold(<.3357). Four rows with blanks in
#familiarity were also deleted.
music <- na.omit(music)
#Copy original data to a new dataframe music1 and exclude unneeded data
music <- music[-c(1:5,7,16,19,21:25,30,34)]
music$artist.hottness.label <- as.factor(music$artist.hottness.label)
str(music)

## 'data.frame': 5648 obs. of 22 variables:
## $ bars_confidence : num 0.643 0.175 0.806 0.873 0.018 0.013 1 0.507 0.125 0.03 ...
## $ beats_confidence : num 0.834 0.883 0.44 0.873 1 0.699 1 0 0.768 1 ...
## $ beats_start : num 0.585 0.136 1.226 0.112 0.429 ...
## $ duration : num 219 210 270 219 245 ...
## $ end_of_fade_in : num 0.247 0.066 5.3 2.125 0.357 ...
## $ familiarity : num 0.582 0.651 0.427 0.36 0.545 ...
## $ key : num 1 2 4 5 7 9 10 7 8 7 ...
## $ key_confidence : num 0.736 0.092 0.717 0.354 0.07 0.205 0 1 0.041 0.725 ...
## $ latitude : num 37.2 37.2 37.2 35.2 37.2 ...
## $ longitude : num -63.9 -63.9 -63.9 -80 -63.9 ...
## $ loudness : num -11.2 -4.5 -13.5 -10.02 -7.54 ...
## $ mode_confidence : num 0.636 0.371 0.652 0.485 0.686 0.305 0.198 0.829 0.516 0.756 ...
## $ start_of_fade_out : num 219 199 259 207 227 ...
## $ tatums_confidence : num 0.779 1 0.121 0.229 0.728 1 0.774 0.377 0.767 0.238 ...

```

```
## $ tatums_start      : num  0.285 0.136 1.226 0.112 0.173 ...
## $ tempo             : num  92.2 129.7 86.6 146.8 118 ...
## $ terms_freq        : num  1 0.887 0.96 0.956 1 ...
## $ time_signature     : num  4 4 4 1 4 4 1 4 5 4 ...
## $ time_signature_confidence: num  0.778 0.562 0.487 0 0.835 0 0.319 0.756 0.579 0.931 ...
## $ year              : int   0 2007 1984 0 0 0 0 1987 0 2004 ...
## $ artist.hotttnesss.label : Factor w/ 3 levels "Cold","Hot","Warm": 3 3 1 1 3 3 1 3 1 2 ...
## $ artist.hotness.label   : Factor w/ 5 levels "Cool","Frigid",...: 4 4 1 2 1 1 2 4 1 5 ...
```

Features


```
#View the number of Cold/Warm/Hot labels
table(music$artist.hotttnesss.label)
```

```
##
## Cold Hot Warm
## 1180 1579 2889
```

```
#View the number of Frigid/Cool/Tepid/Warm/Hot labels
table(music$artist.hotness.label)
```

```
##
## Cool Frigid Hot Tepid Warm
## 1444 973 278 1566 1387
```

```
#Plot artists latitude and longitude
plot(music$latitude,music$longitude)
```



```
#Plot artist hotttnesss
#hist(music$artist.hotttnesss,breaks=20)
#hist(music$artist.hotness,breaks=20)
```

```
#THIS IS INCOMPLETE CODE FOR PLOTTING ADDITIONAL DATA... #Create a map of the world
mapWorld <- borders("world", colour="gray50", fill="white")
```

```

#Code from https://rpubs.com/spoonerf/global_map #Need to figure out what to put in locs
locs<-read.csv("my_locations.csv") locs<- sp_dups<-data.frame(ddply(locs,.(Longitude,Latitude),nrow))
sp_dupsloc_id <- 1 : length(sp_dupsLongitude) sp_dups_df<-merge(sp_dups, locs, by=c("Longitude","Latitude"))

loc<-data.frame(sp_dups_dfLongitude, sp_dups_dLatitude,sp_dups_df$V1) loc<-unique(loc) colnames(loc)<-
c("Longitude", "Latitude", "V1")

coordinates(loc)<-c("Longitude","Latitude") proj4string(loc) <- CRS("+proj=longlat")

loc_df<-data.frame(loc)

theme_opts <- list(theme(panel.grid.minor = element_blank(), panel.grid.major = element_blank(),
panel.background = element_blank(), plot.background = element_rect(fill="white"), panel.border =
element_blank(), axis.line = element_blank(), axis.text.x = element_blank(), axis.text.y = element_blank(),
axis.ticks = element_blank(), axis.title.x = element_blank(), axis.title.y = element_blank(), plot.title =
element_text(size=22)))

library(maps) library(mapdata)

ggplot(data=loc_df, aes(Longitude, Latitude, group=NULL,fill=NULL,size=V1))+#, fill=hole)) + bor-
ders(fill="light grey",colour="light grey")+ geom_point(color="black",alpha=I(7/10))+ scale_size(range=c(1,7),
guide = "legend",labs(size="No. of Populations"))+ coord_equal()+ theme_opts

```

Methods

```

#Do analysis to determine hot/warm/cold artists based on hotttnesss
#The random forest analysis is from a training video by Bharatendra Rai
#at https://www.youtube.com/watch?v=dJclNIN-TPo
#Data Partition - ind = independent samples
#The code below runs in console but not R Markdown
set.seed(123)
ind<- sample(2,nrow(music), replace=TRUE,prob=c(0.7,0.3))
train <- music[ind==1,]
test <- music[ind==2,]
#Run randomForest on 3 levels
library(randomForest)

## randomForest 4.6-14

## Type rfNews() to see new features/changes/bug fixes.

##
## Attaching package: 'randomForest'

## The following object is masked from 'package:ggplot2':
##
##      margin

set.seed(222)
rf <- randomForest(music[, -21:-22],music[,21])
print(rf)

##
## Call:
## randomForest(x = music[, -21:-22], y = music[, 21])
##              Type of random forest: classification
##              Number of trees: 500
## No. of variables tried at each split: 4

```

```
##
##          OOB estimate of  error rate: 20.18%
## Confusion matrix:
##          Cold  Hot  Warm  class.error
## Cold   728    4  448    0.3830508
## Hot     6 1289   284    0.1836605
## Warm   200   198 2491    0.1377639
```

```
attributes(rf)
```

```
## $names
## [1] "call"          "type"          "predicted"
## [4] "err.rate"      "confusion"     "votes"
## [7] "oob.times"     "classes"       "importance"
## [10] "importanceSD"  "localImportance" "proximity"
## [13] "ntree"         "mtry"          "forest"
## [16] "y"            "test"          "inbag"
##
## $class
## [1] "randomForest"
```

```
rf$confusion
```

```
##          Cold  Hot  Warm  class.error
## Cold   728    4  448    0.3830508
## Hot     6 1289   284    0.1836605
## Warm   200   198 2491    0.1377639
```

```
#Run randomForest on 5 levels
```

```
library(randomForest)
set.seed(222)
rf2 <- randomForest(music[, -21:-22], music[, 22])
print(rf2)
```

```
##
## Call:
## randomForest(x = music[, -21:-22], y = music[, 22])
##          Type of random forest: classification
##          Number of trees: 500
## No. of variables tried at each split: 4
##
##          OOB estimate of  error rate: 35.64%
## Confusion matrix:
##          Cool  Frigid  Hot  Tepid  Warm  class.error
## Cool    845    175    0   387    37    0.4148199
## Frigid   279    620    0    73     1    0.3627955
## Hot        1      2   56    15   204    0.7985612
## Tepid    339     45    0   982   200    0.3729246
## Warm      21      6   10   218 1132    0.1838500
```

```
attributes(rf2)
```

```
## $names
## [1] "call"          "type"          "predicted"
## [4] "err.rate"      "confusion"     "votes"
## [7] "oob.times"     "classes"       "importance"
## [10] "importanceSD"  "localImportance" "proximity"
```


```
## [13] "ntree"          "mtry"          "forest"
## [16] "y"             "test"          "inbag"
##
## $class
## [1] "randomForest"
```

```
rf2$confusion
```

```
##      Cool Frigid Hot Tepid Warm class.error
## Cool   845   175  0   387   37   0.4148199
## Frigid  279   620  0    73    1   0.3627955
## Hot      1     2  56   15  204   0.7985612
## Tepid   339    45  0   982  200   0.3729246
## Warm    21     6  10   218 1132   0.1838500
```


```
#Run randomForest again with tune mtry data from below
#rfx <- randomForest(artist.hotness.label ~.,data=music,ntree=200,mtry=8,importance=TRUE,proximity=TRUE)
#Prediction & Confusion Matrix - train data
#library(caret)
#p1<-predict(rfx,train)
#confusionMatrix(p1,train)
#Prediction & Confusion Matrix - test data
#p2<-predict(rfx,test)
#confusionMatrix(p2,test$artist.hotness.label)
#Error rate of Random Forest
plot(rf)
```

rf



```
plot(rf2)
```

rf2



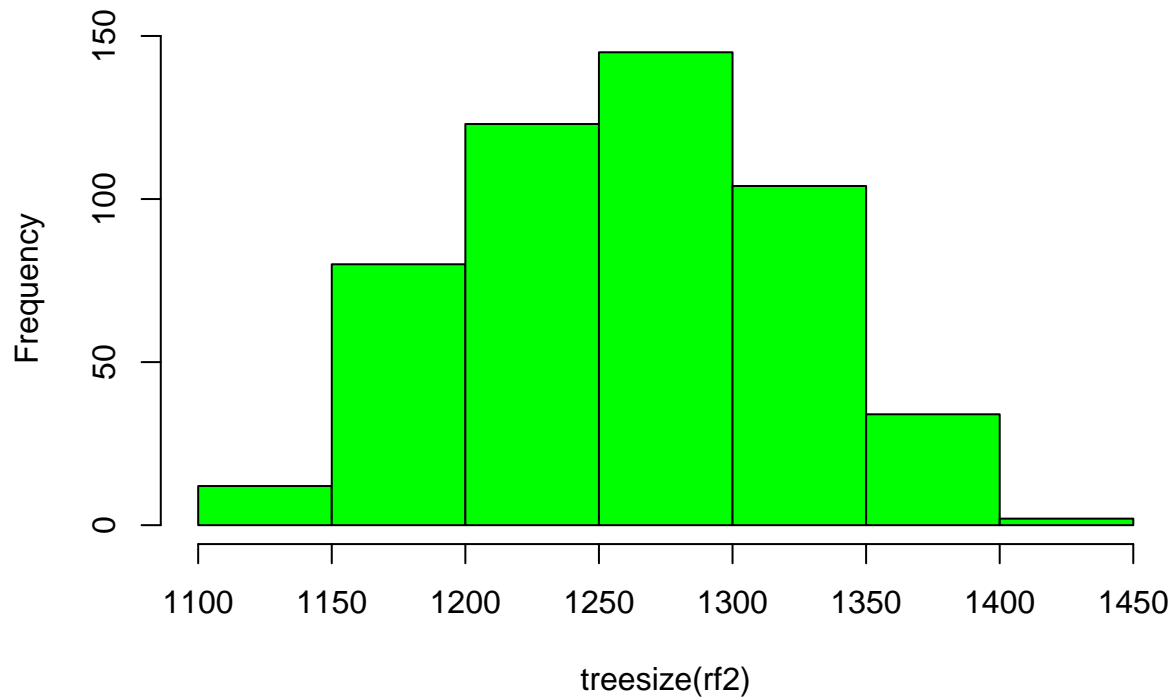
```
#The error rate is not improving after ~100 trees
#Tune mtry
#t <- tuneRF(train[,-21],train[,21],
#           stepFactor=.5,
#           plot=TRUE,
#           ntreeTry=200,
#           trace=TRUE,
#           improve=0.05)
#No. of nodes for the trees
hist(treesize(rf),
     main="Number of Nodes for the Trees",
     col="green")
```


Number of Nodes for the Trees




```
hist(treesize(rf2),  
     main="Number of Nodes for the Trees",  
     col="green")
```

Number of Nodes for the Trees



```
# Variable Importance
# Familiarity is much more important than the other variables. Should it be removed and run again?
varImpPlot(rf,
            sort=T,
            n.var=10,
            main="Top 10 - Variable Importance")
```

Top 10 – Variable Importance



```
importance(rf)
```


##	MeanDecreaseGini
## bars_confidence	107.58378
## beats_confidence	94.47263
## beats_start	114.60290
## duration	111.20674
## end_of_fade_in	88.18414
## familiarity	1520.73102
## key	70.97697
## key_confidence	107.84287
## latitude	80.31500
## longitude	80.00661
## loudness	160.85558
## mode_confidence	112.40895
## start_of_fade_out	109.94686
## tatums_confidence	100.29904
## tatums_start	113.12138
## tempo	118.86929
## terms_freq	142.67373
## time_signature	31.82993
## time_signature_confidence	81.87304
## year	133.56141

```
varUsed(rf)
```

```
## [1] 22981 20248 24173 23234 18362 42545 16709 23076 13678 13543 26260  
## [12] 23733 23098 21580 23800 24889 17934 7870 18132 13527
```

```
varImpPlot(rf2,  
            sort=T,  
            n.var=10,  
            main="Top 10 - Variable Importance")
```

Top 10 – Variable Importance



```
importance(rf2)
```

```
##                               MeanDecreaseGini  
## bars_confidence                174.77925  
## beats_confidence               153.37570  
## beats_start                    176.89591  
## duration                      168.58389  
## end_of_fade_in                 141.16350  
## familiarity                    1362.55235  
## key                           115.59861  
## key_confidence                 167.04577  
## latitude                      114.97836  
## longitude                     118.58999  
## loudness                      227.38098  
## mode_confidence               176.77603  
## start_of_fade_out              167.99820  
## tatums_confidence             160.88208  
## tatums_start                  178.93330  
## tempo                         186.65606  
## terms_freq                    176.06618  
## time_signature                 51.64041
```

```
## time_signature_confidence      134.38659
## year                           167.90690
varUsed(rf2)

## [1] 35548 31313 36142 34492 28373 59265 26210 34634 19869 20060 39304
## [12] 36051 34736 33307 36125 37324 26336 12053 28255 20695

#Multidimensional Scaling Plot
#The code below causes R to lock up...
#MDSplot(rf,train$artist.hottnesss.label)
```

Results

Conclusion

Appendices